


AWS interview questions along with their answers

Follow for More – Krishan Bhatt 

1. What is EC2? How does it differ from traditional hosting?

- Answer: Amazon Elastic Compute Cloud (EC2) is a web service that provides resizable compute capacity in the cloud. It allows users to rent virtual servers, known as instances, and run applications on them. EC2 differs from traditional hosting in that it offers scalability, flexibility, and cost-effectiveness, as users only pay for the resources they use, without the need for long-term contracts or upfront payments.

2. Explain the difference between S3 and EBS.

- Answer: Amazon Simple Storage Service (S3) is an object storage service that provides scalable storage for data objects, while Amazon Elastic Block Store (EBS) is a block storage service that provides persistent block-level storage volumes for use with EC2 instances. S3 is ideal for storing large amounts of unstructured data, such as photos, videos, and backups, while EBS volumes are used as storage drives for EC2 instances, providing low-latency access to data.

3. What is IAM? How do you manage access in AWS?

- Answer: AWS Identity and Access Management (IAM) is a web service that helps securely control access to AWS resources. It allows users to create and manage AWS users and groups, and assign permissions to them to access or manage AWS resources. Access in AWS is managed through IAM policies, which define the actions users are allowed or denied to perform on AWS resources.

4. What is a VPC? How do you create one?

- Answer: A Virtual Private Cloud (VPC) is a virtual network dedicated to a user's AWS account. It allows users to define a virtual network topology, including subnets, route tables, and security groups, and launch AWS resources, such as EC2 instances and RDS databases, within the VPC. A VPC can be created using the AWS Management Console, AWS CLI, or AWS SDKs by specifying the desired IP address range, subnets, and other network configuration settings.

5. How do you secure data at rest and in transit in AWS?

- Answer: Data at rest can be secured in AWS using encryption techniques, such as AWS Key Management Service (KMS) for managing encryption keys and encrypting data stored in S3, EBS, and RDS. Data in transit can be secured using Secure Sockets Layer (SSL)/Transport Layer Security (TLS) protocols for encrypting data transmitted over the network, and Virtual Private Network (VPN) connections or AWS

Direct Connect for establishing secure connections between on-premises data centers and AWS.

6. What is CloudWatch? How do you use it for monitoring?

- Answer: Amazon CloudWatch is a monitoring and observability service that provides real-time monitoring of AWS resources and applications. It collects and tracks metrics, logs, and events, allowing users to gain insights into resource utilization, application performance, and operational health. CloudWatch can be used to set alarms, visualize metrics, and automate actions based on predefined thresholds.

7. Explain the purpose of CloudTrail.

- Answer: AWS CloudTrail is a service that enables governance, compliance, operational auditing, and risk auditing of AWS accounts. It records API calls and actions taken by users, applications, or AWS services within an AWS account, providing a history of activity for security analysis, resource change tracking, and troubleshooting. CloudTrail logs can be stored in Amazon S3 and analyzed using Amazon Athena, Amazon Elasticsearch Service, or third-party log analysis tools.

8. How do you design for high availability in AWS?

- Answer: Designing for high availability in AWS involves implementing redundancy and fault tolerance across multiple availability zones (AZs) within a region. This includes deploying resources across multiple AZs, using load balancers for distributing traffic, implementing auto-scaling for handling sudden increases in demand, and leveraging AWS services like Amazon Route 53 for DNS routing and failover.

9. What is Auto Scaling? How does it work?

- Answer: Auto Scaling is a feature of AWS that automatically adjusts the number of EC2 instances in a fleet based on demand or predefined metrics. It helps maintain application availability, optimize performance, and reduce costs by dynamically scaling resources up or down in response to changes in workload. Auto Scaling can be configured using scaling policies and triggers, which define when and how instances are launched or terminated.

10. What is RDS? When would you use it over self-managed databases?

- Answer: Amazon Relational Database Service (RDS) is a managed database service that simplifies the deployment, management, and scaling of relational databases in the cloud. RDS supports several database engines, including MySQL, PostgreSQL, MariaDB, Oracle, and SQL Server. It is preferred over self-managed databases when organizations want to offload routine database administration tasks,

such as backups, patching, and scaling, to AWS, allowing them to focus on application development and business logic.

11. What is AWS Lambda? How does it work?

- Answer: AWS Lambda is a serverless computing service that allows users to run code without provisioning or managing servers. It automatically scales to handle incoming requests and charges only for the compute time consumed. Developers can upload code in the form of functions to Lambda, which executes the code in response to events, such as HTTP requests, changes to data in S3, or messages from Amazon SQS.

12. How do you deploy a serverless application in AWS?

- Answer: Serverless applications in AWS are typically deployed using AWS Serverless Application Model (SAM) or AWS CloudFormation templates. Developers define the application's resources, including Lambda functions, API Gateway endpoints, and other AWS services, in a template file, which is then deployed using the AWS CLI or AWS Management Console. Alternatively, serverless frameworks like Serverless Framework or AWS Amplify can be used for deployment.

13. How do you optimize costs in AWS?

- Answer: Cost optimization in AWS involves analyzing resource usage, identifying opportunities for savings, and implementing cost-saving strategies. This can include rightsizing EC2 instances, using Reserved Instances or Savings Plans for predictable workloads, leveraging spot instances for flexible workloads, implementing auto-scaling to match resource usage with demand, and using AWS Cost Explorer for cost analysis and forecasting.

14. What are some best practices for reducing AWS bills?

- Answer: Some best practices for reducing AWS bills include enabling cost allocation tags for tracking resource usage, leveraging AWS Trusted Advisor for cost optimization recommendations, implementing lifecycle policies for managing data stored in S3, optimizing data transfer costs by using AWS Direct Connect or Amazon CloudFront, and using AWS Budgets for setting and monitoring cost budgets.

15. How do you manage access control in AWS using IAM?

- Answer: Access control in AWS is managed using IAM, which allows users to create and manage IAM users, groups, roles, and policies. IAM policies define permissions that specify what actions users are allowed or denied to perform on AWS resources. Access can be granted based on principles of least privilege, using IAM roles for temporary access, and implementing multi-factor authentication (MFA) for added security.

16. What is AWS CloudFormation? How do you use it?

- Answer: AWS CloudFormation is a service that allows you to provision and manage AWS infrastructure as code using templates. These templates are written in JSON or YAML and define the resources and their configurations needed for an application. CloudFormation automates the deployment process, making it easier to manage and replicate infrastructure across environments.

17. Explain the difference between horizontal scaling and vertical scaling. When would you use each in AWS?

- Answer: Horizontal scaling involves adding more instances to distribute the load across multiple machines, while vertical scaling involves increasing the resources (CPU, memory, etc.) of existing instances. Horizontal scaling is typically used for stateless applications that can easily scale out, while vertical scaling is used for stateful applications that require more resources on a single instance.

18. What is AWS Elastic Beanstalk? How does it work?

- Answer: AWS Elastic Beanstalk is a Platform as a Service (PaaS) offering that makes it easy to deploy and manage web applications and services. It automatically handles the deployment, capacity provisioning, load balancing, scaling, and monitoring of the application, allowing developers to focus on writing code. Elastic Beanstalk supports various programming languages and frameworks, including Java, .NET, Node.js, Python, and Docker.

19. How do you ensure data durability and availability in AWS S3?

- Answer: AWS S3 provides eleven 9's (99.999999999%) of durability by automatically replicating data across multiple facilities within a region. To ensure availability, users can enable versioning to protect against accidental deletion or overwrite, configure cross-region replication for disaster recovery, and use S3 Transfer Acceleration to improve data transfer speeds.

20. Explain the concept of a Content Delivery Network (CDN) and how it's used in AWS.

- Answer: A Content Delivery Network (CDN) is a distributed network of servers that delivers web content to users based on their geographic location, reducing latency and improving performance. In AWS, Amazon CloudFront is a CDN service that caches content at edge locations around the world, providing low-latency access to static and dynamic content, such as images, videos, and APIs, from the nearest edge location to the end user.

21. What is Amazon RDS Multi-AZ deployment?

- Answer: Amazon RDS Multi-AZ (Multi-Availability Zone) deployment is a high availability feature that automatically replicates your database to a standby instance in a different Availability Zone (AZ). In the event of a planned or unplanned outage, Amazon RDS automatically fails over to the standby instance, ensuring database availability and data durability.

22. Explain the difference between Amazon DynamoDB and Amazon RDS.

- Answer: Amazon DynamoDB is a fully managed NoSQL database service that provides fast and predictable performance at any scale, while Amazon RDS is a managed relational database service that supports multiple database engines (such as MySQL, PostgreSQL, SQL Server, etc.). DynamoDB is ideal for applications requiring low-latency, high-throughput, and flexible scaling, while RDS is suitable for traditional relational database workloads with structured data requirements.

23. What is AWS Lambda Layers? How do you use them?

- Answer: AWS Lambda Layers allow you to centrally manage code and data that you want to reuse across multiple Lambda functions. Layers can contain libraries, custom runtimes, or other dependencies that are shared across functions. By using layers, you can reduce duplication of code and simplify maintenance across your serverless applications.

24. What is AWS CloudTrail Insights?

- Answer: AWS CloudTrail Insights is a feature of AWS CloudTrail that uses machine learning algorithms to analyze API activity logs and detect unusual or suspicious behavior. It automatically identifies and alerts users to potentially risky actions, unauthorized access attempts, or security threats, helping improve security posture and compliance in AWS environments.

25. Explain the difference between Amazon S3 Standard, S3 Standard-IA, and S3 Glacier storage classes.

- Answer: Amazon S3 offers several storage classes designed for different use cases and access patterns. S3 Standard is the default storage class for frequently accessed data with low latency requirements. S3 Standard-IA (Infrequent Access) is for data accessed less frequently but requires rapid access when needed. S3 Glacier is for long-term archival and backup data with retrieval times ranging from minutes to hours. Each storage class offers different pricing and durability options to optimize costs based on data access patterns.

26. What is AWS Elastic Load Balancing? How does it work?

- Answer: AWS Elastic Load Balancing (ELB) automatically distributes incoming application traffic across multiple targets, such as EC2 instances, containers, IP addresses, or Lambda functions, to ensure high availability and fault tolerance of your application. ELB operates at the transport layer (Layer 4) with Network Load Balancers (NLB) and at the application layer (Layer 7) with Application Load Balancers (ALB) and Classic Load Balancers.

27. Explain the concept of Cross-Origin Resource Sharing (CORS) and how it's implemented in AWS.

- Answer: Cross-Origin Resource Sharing (CORS) is a mechanism that allows web browsers to make cross-origin requests securely. In AWS, CORS can be configured for Amazon S3 buckets to specify which origins are allowed to access the bucket's resources using HTTP headers. By configuring CORS rules, you can control access to resources stored in S3 buckets from different origins (domains).

28. What is AWS Direct Connect?

- Answer: AWS Direct Connect is a network service that provides a dedicated private connection between your on-premises data center or corporate network and AWS. It bypasses the public internet, offering lower latency, increased security, and consistent network performance for accessing AWS resources. Direct Connect is ideal for applications requiring high bandwidth, low latency, and secure connectivity to AWS services.

29. Explain the concept of AWS Virtual Private Cloud (VPC) Peering.

- Answer: AWS VPC Peering allows you to connect two VPCs within the same AWS region, enabling communication between instances in the peered VPCs using private IP addresses. Peering connections are secure and private, and they can be established between VPCs belonging to the same AWS account or different AWS

accounts. VPC peering simplifies network architecture and allows you to extend your network across multiple VPCs.

30. What is AWS Snowball?

- Answer: AWS Snowball is a petabyte-scale data transfer service that allows you to securely transfer large amounts of data into and out of AWS. It consists of ruggedized storage devices, called Snowball devices, that are shipped to your location, where you can transfer data onto them using standard NAS or SAN protocols. Once the data transfer is complete, the Snowball device is returned to AWS for data ingestion into your AWS environment. Snowball is particularly useful for data migration, data backup, or large-scale data transfer projects with limited internet bandwidth.

31. What is AWS Key Management Service (KMS)?

- Answer: AWS Key Management Service (KMS) is a managed service that allows you to create and control encryption keys used to encrypt your data. KMS makes it easy to create and manage cryptographic keys and control their use across a wide range of AWS services and in your applications. It integrates with AWS CloudTrail for auditing key usage and offers hardware security module (HSM) protection for sensitive key material.

32. Explain the concept of AWS Identity Federation.

- Answer: AWS Identity Federation allows you to grant temporary access to AWS resources to users who authenticate outside of AWS, such as employees or partners who use their existing corporate credentials. It enables single sign-on (SSO) for users across multiple systems and applications by federating identity information from an external identity provider (IdP), such as Active Directory, LDAP, or SAML-based identity providers.

33. What is AWS CodePipeline? How does it work?

- Answer: AWS CodePipeline is a continuous integration and continuous delivery (CI/CD) service that automates the build, test, and deployment phases of your software release process. It allows you to define a series of stages and actions in a pipeline that automatically trigger when changes are made to your source code repository. CodePipeline integrates with other AWS services, such as AWS CodeBuild, AWS CodeDeploy, and AWS CodeCommit, to orchestrate and automate the entire release process.

34. Explain the difference between AWS ECS and AWS EKS.

- Answer: AWS ECS (Elastic Container Service) is a fully managed container orchestration service for running Docker containers on EC2 instances. It simplifies the deployment and management of containerized applications at scale. AWS EKS (Elastic Kubernetes Service) is a fully managed Kubernetes service that allows you to run Kubernetes clusters on AWS without the need to install, operate, or maintain the underlying infrastructure. ECS is ideal for users who prefer a simpler and more integrated container solution, while EKS is suitable for users who require the flexibility and extensibility of Kubernetes.

35. What is AWS DataSync?

- Answer: AWS DataSync is a data transfer service that makes it easy to automate and accelerate data movement between on-premises storage systems and AWS storage services, such as Amazon S3, Amazon EFS, and Amazon FSx for Windows File Server. It uses a purpose-built protocol optimized for high-speed, secure data transfer and can handle large-scale data migration, replication, and synchronization tasks with minimal setup and management overhead.●

36. What is AWS CloudWatch Logs? How do you use it?

- Answer: AWS CloudWatch Logs is a monitoring and log management service that allows you to collect, view, and analyze log data generated by your AWS resources and applications. You can use CloudWatch Logs to centralize logs from

multiple sources, create metric filters and alarms based on log events, and troubleshoot issues by searching and filtering log data.

37. Explain the difference between Amazon S3 event notifications and Amazon S3 inventory.

- Answer: Amazon S3 event notifications allow you to receive notifications when certain events occur in your S3 buckets, such as object creation, deletion, or replication. S3 event notifications can trigger actions in other AWS services, such as invoking AWS Lambda functions or sending messages to Amazon SNS topics. On the other hand, Amazon S3 inventory provides a scheduled report of your bucket's metadata, including information about all objects stored in the bucket, which can be used for auditing, compliance, and analytics purposes.

38. What is AWS CloudFormation change sets?

- Answer: AWS CloudFormation change sets allow you to preview and manage changes to your AWS infrastructure before applying them. You can create a change set to review the proposed changes to your CloudFormation stack, including additions, modifications, and deletions of resources, and then execute or discard the change set based on your review. Change sets help ensure that infrastructure changes are applied safely and predictably.

39. How do you implement data encryption in transit and at rest in AWS RDS?

- Answer: In AWS RDS, data encryption in transit can be implemented by enabling SSL/TLS encryption for connections to the database instance, which encrypts data as it travels between the client application and the database instance over the network. Data encryption at rest can be achieved by enabling the encryption option when creating or modifying the RDS instance, which encrypts the database storage volumes using AWS Key Management Service (KMS) keys.

40. What is AWS Lambda@Edge?

- Answer: AWS Lambda@Edge is an extension of AWS Lambda that allows you to run serverless functions at edge locations of the AWS global network, closer to your end users. Lambda@Edge functions can intercept and modify requests and responses at the edge, enabling you to customize content delivery, implement security measures, and optimize performance for your web applications and content delivery network (CDN) distributions.

41. Explain the difference between Amazon DynamoDB Streams and Amazon Kinesis Streams.

- Answer: Amazon DynamoDB Streams is a feature of Amazon DynamoDB that captures a time-ordered sequence of item-level modifications made to DynamoDB tables and makes these changes available for consumption in near real-time. DynamoDB Streams is primarily used for building applications that react to changes in DynamoDB data, such as data replication, indexing, or stream processing. Amazon Kinesis Streams, on the other hand, is a fully managed stream processing service that allows you to collect and process large streams of data records in real-time. Kinesis Streams is designed for scenarios that require real-time analytics, data processing, and event-driven applications.

42. What is AWS Storage Gateway?

- Answer: AWS Storage Gateway is a hybrid cloud storage service that enables on-premises applications to seamlessly access data stored in AWS cloud storage services, such as Amazon S3, Amazon Glacier, and Amazon EBS. Storage Gateway provides different types of gateways to meet various use cases, including file gateway for file-based access, volume gateway for block-based access, and tape gateway for virtual tape library (VTL) backup and archival.

43. Explain the AWS Well-Architected Framework and its pillars.

- Answer: The AWS Well-Architected Framework is a set of best practices and guidelines for building secure, high-performing, resilient, and efficient cloud architectures. It consists of five pillars: Operational Excellence, Security, Reliability, Performance Efficiency, and Cost Optimization. Each pillar provides principles, design patterns, and recommendations for designing and evaluating cloud architectures to ensure they meet the requirements of modern cloud-native applications.

44. What is Amazon Redshift? How does it differ from traditional relational databases?

- Answer: Amazon Redshift is a fully managed data warehouse service that allows you to analyze large datasets using SQL queries. It is optimized for online analytical processing (OLAP) workloads and provides high-performance, scalable, and cost-effective data warehousing solutions in the cloud. Redshift differs from traditional relational databases in terms of its distributed and columnar storage architecture, massively parallel processing (MPP) capabilities, and integration with other AWS services for data ingestion, transformation, and visualization.

45. How do you monitor and manage costs in AWS using AWS Budgets?

- Answer: AWS Budgets is a service that allows you to set custom budgets for monitoring and controlling your AWS spending. You can create budgets based on cost, usage, or reservation utilization, and receive alerts when your actual spending exceeds or is forecasted to exceed your budgeted amount. AWS Budgets helps you track your AWS spending, identify cost trends, and take proactive actions to optimize costs and resource utilization.

46. What is Amazon Elastic File System (EFS)? How does it differ from Amazon EBS?

- Answer: Amazon Elastic File System (EFS) is a fully managed file storage service that provides scalable and highly available network file systems that can be shared across multiple EC2 instances. EFS is designed for workloads that require shared access to a common data source, such as content management systems, web serving, and development environments. EFS differs from Amazon Elastic Block Store (EBS) in that EFS provides file-level storage, whereas EBS provides block-level storage volumes that are attached to individual EC2 instances. EFS also supports scalable capacity and automatic scaling, whereas EBS volumes have fixed size and performance characteristics.

47. What is AWS Step Functions? How does it work?

- Answer: AWS Step Functions is a fully managed service that allows you to coordinate and orchestrate distributed applications and microservices using visual workflows. Step Functions enables you to define complex state machines that specify the sequence of steps, conditions, and error handling logic for executing business processes or workflows. You can trigger Step Functions workflows using events from AWS services, such as Lambda, SQS, SNS, and more, and monitor the execution and status of workflows in real-time using the AWS Management Console or APIs.

48. Explain the concept of AWS Organizations and how it's used for managing multiple AWS accounts.

- Answer: AWS Organizations is a service that allows you to centrally manage and govern multiple AWS accounts within your organization. Organizations provides features for consolidating billing and cost management, enforcing security and compliance policies, and automating account provisioning and management tasks. With Organizations, you can create a hierarchy of organizational units (OUs), apply policies at different levels of the hierarchy, and delegate administrative access to specific accounts or groups of accounts.

49. What is AWS Transit Gateway?

- Answer: AWS Transit Gateway is a fully managed service that simplifies the management and routing of network traffic between multiple VPCs, VPN connections, and on-premises networks. Transit Gateway acts as a central hub that connects multiple VPCs and on-premises networks, allowing you to scale your network architecture and simplify connectivity for hybrid cloud.

50. Explain the difference between AWS EC2 Auto Scaling and AWS ECS Auto Scaling.

- Answer: AWS EC2 Auto Scaling is a feature that automatically adjusts the number of EC2 instances in a fleet based on demand or predefined metrics, ensuring that you have the right amount of compute capacity to handle varying workloads. On the other hand, AWS ECS Auto Scaling is a feature that automatically adjusts the number of tasks or containers in an ECS service based on metrics such as CPU utilization or memory usage, allowing you to scale containerized applications dynamically. While EC2 Auto Scaling operates at the instance level, ECS Auto Scaling operates at the container level within ECS services.